



# **NAVIGATING FUTURE OF ONLINE SHOPPING**

SQL and Python Approach

PRESENTED BY MOHI GUPTA





# DATASET OVERVIEW

01

CUSTOMERS.CSV:  
CUSTOMERS  
DEMOGRAPHIC

02

SELLERS.CSV : SELLERS  
INFORMATION

03

ORDER\_ITEMS : ORDER  
ITEMS DETAILS

04

GEOLOCATION.CSV :  
GEOLOCATION DETAILS

05

PAYMENTS.CSV :  
PAYMENTS DETAILS

06

ORDERS.CSV : ORDER  
ISTORY AND DETAILS

07

PRODUCTS.CSV :  
PRODUCTS DETAILS

Goal: Analyze and visualize business metrics using SQL and Python



# Problem Statement

In the fast-growing e-commerce and retail industry, understanding customer behavior and sales patterns is crucial for optimizing business operations and improving customer retention. This project aims to extract and analyze key business metrics from an e-commerce dataset, which includes customer orders, payments, products, and other transactional details.

- The primary challenge lies in deriving actionable insights such as:
- How do customers behave over time in terms of spending and order frequency?
- What trends can be identified for year-over-year sales growth?
- How can customer retention be measured and improved?
- Which customers contribute the most to the total revenue, and how can they be targeted for personalized marketing?

# Objective

THE OBJECTIVE OF THIS PROJECT IS TO:

- Analyze the dataset using SQL for complex queries and Python for data manipulation and visualization.
- Identify strategic insights like customer retention rates, moving averages, and year-over-year growth in total sales.
- Generate customer-centric insights, such as identifying top spenders and calculating retention rates based on repeated purchases within a specific time frame.
- Create visualizations to communicate trends, correlations, and key findings clearly.

# Basic Problems

OBJECTIVE: EXTRACT FUNDAMENTAL INSIGHTS FROM THE DATASET

- List all unique cities where customers are located.

```
SELECT DISTINCT customer_city  
FROM customers  
WHERE customer_city IS NOT NULL  
ORDER BY customer_city;
```

customer_city
abadia dos dourados
abadiania
abaete
abaetetuba
abaiara
abaira
abare
abatia
abdon batista
abelardo luz
abranes
abre campo
abreu e lima

```
▶ customer_cities = (  
    final_df['customer_city']  
    .dropna()  
    .str.strip()  
    .str.lower()  
    .unique()  
customer_cities = sorted(customer_cities)  
for city in customer_cities:  
    print(city)
```

```
⇒ abaete  
abaetetuba  
abaiara  
abelardo luz  
abranes  
abre campo  
abreu e lima  
acaiaca  
acailandia  
acopiara  
acreuna  
acucena  
adamantina  
adolfo  
adustina  
afogados da ingazeira  
afonso claudio  
afranio
```



- COUNT THE NUMBER OF ORDERS PLACED IN 2017.

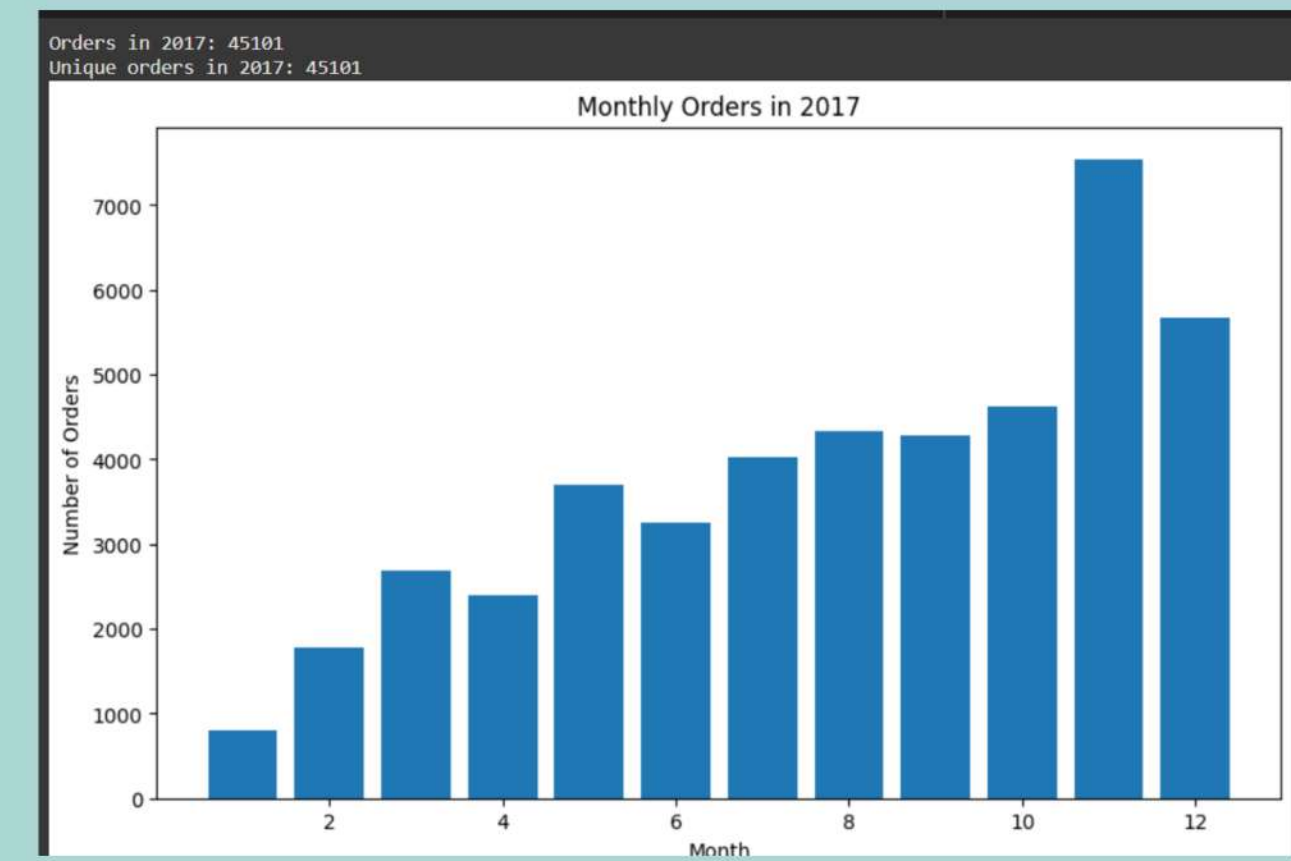
```
SELECT COUNT(*) AS order_count
FROM orders
WHERE YEAR(order_purchase_timestamp) = 2017
      AND order_purchase_timestamp IS NOT NULL;
```

Result Grid		Filter Rows:
	order_count	
▶	45101	

```
import builtins
print = builtins.print

df_orders['order_purchase_timestamp'] = pd.to_datetime(
    df_orders['order_purchase_timestamp'], format='%d-%m-%Y %H:%M', errors='coerce'
)
orders_count = len(orders_2017)
unique_orders_count = orders_2017['order_id'].nunique()
print("Orders in 2017:", orders_count)
print("Unique orders in 2017:", unique_orders_count)

# Plot
orders_2017['month'] = orders_2017['order_purchase_timestamp'].dt.month
monthly_counts = orders_2017.groupby('month')['order_id'].nunique()
plt.figure(figsize=(10, 6))
plt.bar(monthly_counts.index, monthly_counts.values)
plt.xlabel('Month')
plt.ylabel('Number of Orders')
plt.title('Monthly Orders in 2017')
plt.show()
```





- FIND THE TOTAL SALES PER CATEGORY.

```
SELECT
    products.`product category`,
    SUM(order_items.price) AS total_sales
FROM
    order_items
JOIN
    products ON order_items.product_id = products.product_id
GROUP BY
    products.`product category`
ORDER BY
    total_sales DESC;
```

Result Grid			Filter Rows:
	product category	total_sales	
▶	HEALTH BEAUTY	1258681.34	
	Watches present	1205005.68	
	bed table bath	1036988.68	
	sport leisure	988048.97	
	computer accessories	911954.32	
	Furniture Decoration	729762.49	
	Cool Stuff	635290.85	
	housewares	632248.66	
	automotive	592720.11	

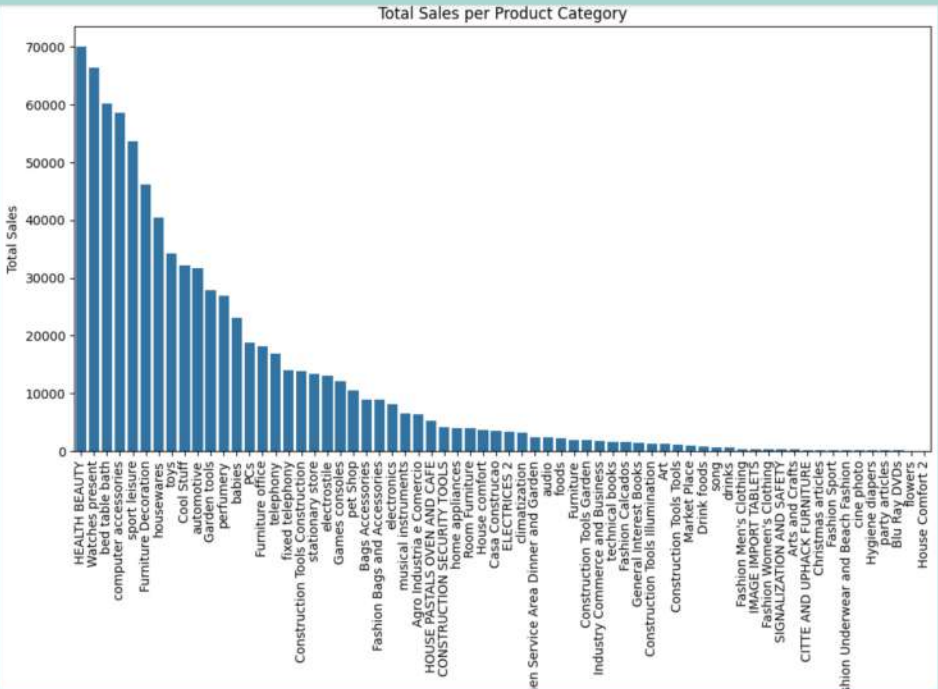
```
[ ] merged = df_order_items.merge(df_products[['product_id', 'product category']],
                                on='product_id', how='left')

# Check for missing categories after merge
print("Missing product categories after merge:", merged['product category'].isnull().sum())

# Group and sum (just like SQL)
total_sales_per_category = merged.groupby('product category')['price'].sum().sort_values(ascending=False)

# Print only categories you see in SQL for visual comparison
print(total_sales_per_category)

#Plot
plt.figure(figsize=(12, 6))
sns.barplot(x=total_sales_per_category.index, y=total_sales_per_category.values)
plt.xlabel('Product Category')
plt.ylabel('Total Sales')
plt.title('Total Sales per Product Category')
plt.xticks(rotation=90)
plt.show()
```



```
Missing product categories after merge: 1603
product category
HEALTH BEAUTY                1258681.34
Watches present              1205005.68
bed table bath               1036988.68
sport leisure                 988048.97
computer accessories         911954.32
...
flowers                       1110.04
House Comfort 2               760.27
cds music dvds               730.00
Fashion Children's Clothing   569.85
insurance and services        283.29
Name: price, Length: 73, dtype: float64
```



- **Calculate the percentage of orders that were paid in installments.**

```
-- 4. Calculate the percentage of orders that were paid in installments.
select 100 * count(case when payment_installments > 1 then 1 end ) / count(*) as percent_installments
from payments;
```

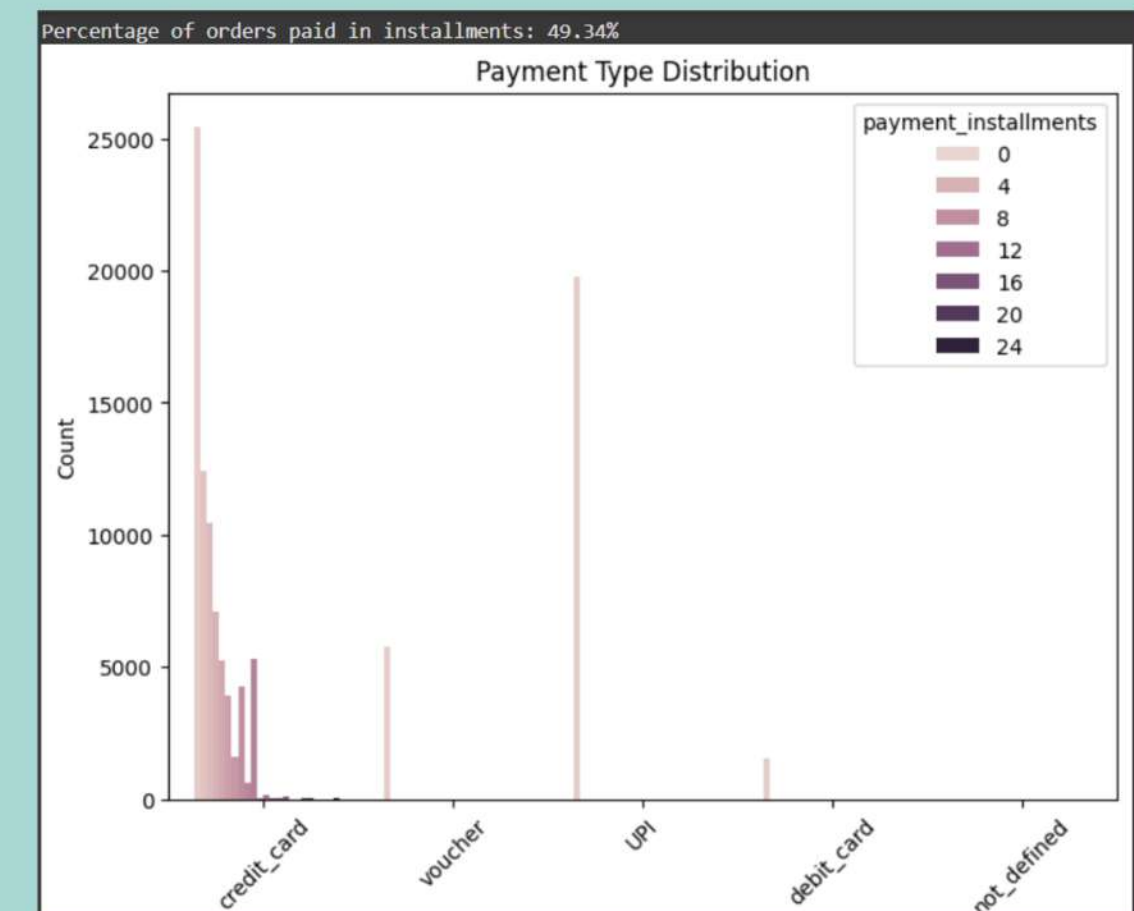
	percent_installments
▶	49.4176

```
# Count the total number of orders
total_orders = orders_payments['order_id'].nunique()
installments_orders = orders_payments[orders_payments['payment_installments'] == 1]

# Count the total number of installments orders
total_installments_orders = installments_orders['order_id'].nunique()

# Calculate the percentage
percentage_installments = (total_installments_orders / total_orders) * 100
print(f"Percentage of orders paid in installments: {percentage_installments:.2f}%")

#Plot
plt.figure(figsize=(8, 6))
sns.countplot(data=orders_payments, x='payment_type', hue='payment_installments')
plt.xlabel('Payment Type')
plt.ylabel('Count')
plt.title('Payment Type Distribution')
plt.xticks(rotation=45)
plt.show()
```





- **COUNT THE NUMBER OF CUSTOMERS FROM EACH STATE.**

```
-- 5. Count the number of customers from each state.
Select customer_state ,
Count(*) AS total_count
from customers group by customer_state;
```

Result Grid			Filter Rows:
	customer_state	total_count	
▶	SP	41746	
	MG	11635	
	ES	2033	
	RJ	12852	
	RS	5466	
	BA	3380	
	CE	1336	
	PR	5045	
	MS	715	

```
▶ total_customers = df_customers['customer_state'].value_counts()
print(total_customers)
#plot
plt.figure(figsize=(12, 6))
sns.barplot(x=total_customers.index, y=total_customers.values)
plt.xlabel('State')
plt.ylabel('Number of Customers')
plt.title('Number of Customers per State')
plt.xticks(rotation=90)
plt.show()
```

```
customer_state
SP      41746
RJ      12852
MG      11635
RS       5466
PR       5045
SC       3637
BA       3380
DF       2140
ES       2033
GO       2020
PE       1652
CE       1336
PA        975
MT        907
MA        747
MS        715
PB        536
```

# INTERMEDIATE PROBLEMS

OBJECTIVE: DIVE DEEPER INTO SALES AND ORDER TRENDS

- Calculate the number of orders per month in 2018

```
-- INTERMEDIATE PROBLEMS
-- Objective: Dive deeper into sales and order trends.
-- 1. Calculate the number of orders per month in 2018.
SELECT MONTH(order_purchase_timestamp) as month ,
count(*) as order_count
from orders
where year(order_purchase_timestamp)=2018
group by MONTH(order_purchase_timestamp)
order by month;
```

```
df_orders['order_purchase_timestamp'] = pd.to_datetime(df_orders['order_purchase_timestamp'], errors='coerce')
df_orders = df_orders.dropna(subset=['order_purchase_timestamp']) # Drop NaT

# Filter for 2018
orders_2018 = df_orders[df_orders['order_purchase_timestamp'].dt.year == 2018]

# Group and count
orders_per_month = orders_2018.groupby(orders_2018['order_purchase_timestamp'].dt.month).size().sort_index()
print(orders_per_month)
```

	month	order_count
▶	1	7269
	2	6728
	3	7211
	4	6939
	5	6873
	6	6167
	7	6292
	8	6512
	9	16

```
order_purchase_timestamp
1      7269
2      6728
3      7211
4      6939
5      6873
6      6167
7      6292
8      6512
9         16
10         4
dtype: int64
```



- Find the average number of products per order, grouped by customer city.

```
-- 2.Find the average number of products per order, grouped by customer city.
SELECT
    customers.customer_city,
    AVG(order_count.products_in_order) AS avg_order
FROM
    (
        SELECT
            order_items.order_id,
            COUNT(order_items.product_id) AS products_in_order
        FROM
            order_items
        GROUP BY
            order_items.order_id
    ) AS order_count
JOIN orders ON order_count.order_id = orders.order_id
JOIN customers ON orders.customer_id = customers.customer_id
GROUP BY
    customers.customer_city
ORDER BY
    avg_order;
```

	customer_city	avg_order
▶	alvares machado	1.0000
	santa fe do sul	1.0000
	penaforte	1.0000
	mampituba	1.0000
	pouso novo	1.0000
	nova america da colina	1.0000
	cerro grande do sul	1.0000
	hora	1.0000

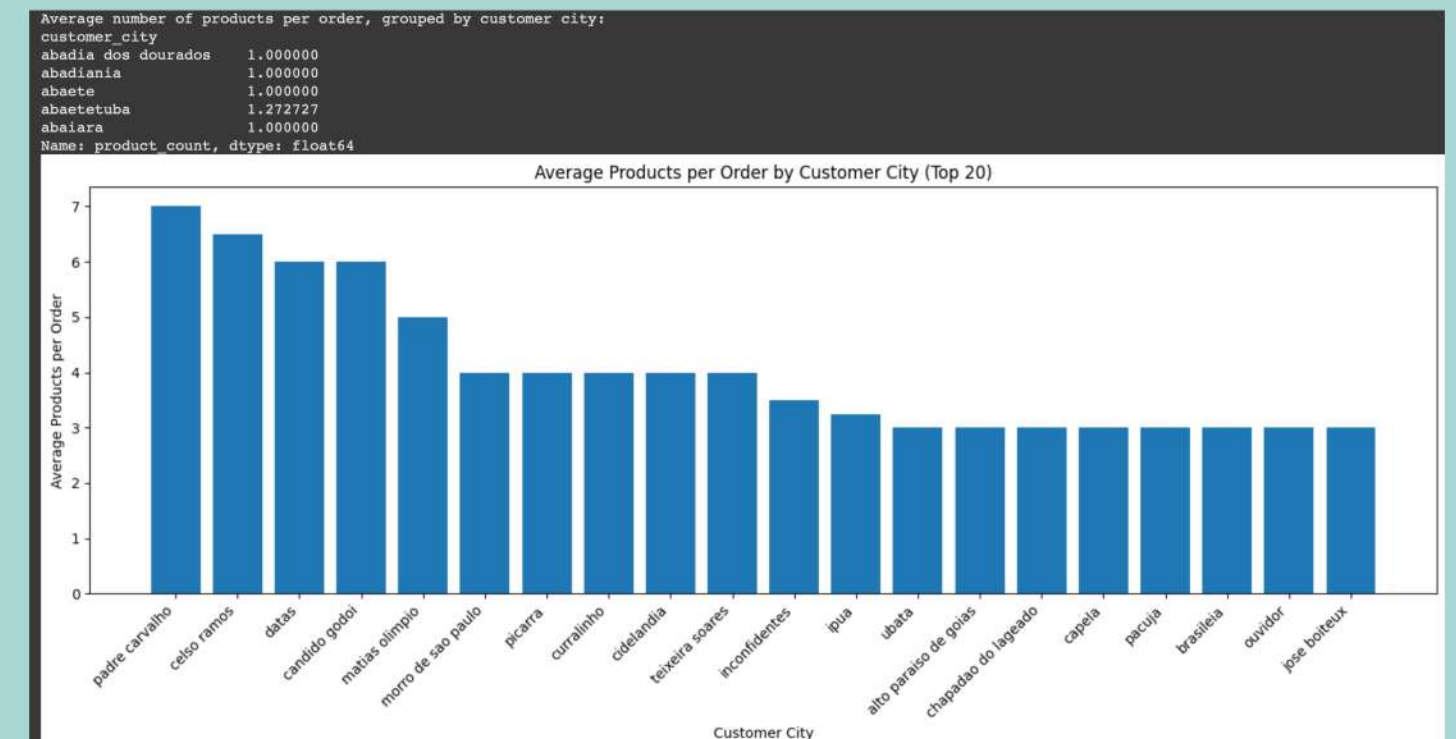
```
products_per_order = df_order_items.groupby('order_id')['product_id'].count().reset_index(name='product_count')

# Now products_per_order is a DataFrame with columns 'order_id' and 'product_count'
products_per_order_per_city = products_per_order.merge(orders_customers[['order_id', 'customer_city']], on='order_id', how='left')

# We now group by 'customer_city' and take the mean of the 'product_count' column
average_products_per_city = products_per_order_per_city.groupby('customer_city')['product_count'].mean()

print("Average number of products per order, grouped by customer city:")
print(average_products_per_city.head())

#plot
top_n = 20
top_cities = average_products_per_city.sort_values(ascending=False).head(top_n)
plt.figure(figsize=(14, 6))
plt.bar(top_cities.index, top_cities.values)
plt.xlabel('Customer City')
plt.ylabel('Average Products per Order')
plt.title('Average Products per Order by Customer City (Top 20)')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```





- Calculate the percentage of total revenue contributed by each product category

```
SELECT
    products.`product category`,
    SUM(order_items.price) AS category_revenue,
    ROUND(100 * SUM(order_items.price) /
        (SELECT SUM(order_items.price)
         FROM order_items
         JOIN products ON order_items.product_id = products.product_id
        ), 2) AS percent_of_total
FROM order_items
JOIN products ON order_items.product_id = products.product_id
GROUP BY products.`product category`
ORDER BY percent_of_total DESC;
```

Result Grid			
Filter Rows:			
Export:			
	product category	category_revenue	percent_of_total
▶	HEALTH BEAUTY	1258681.34	9.26
	Watches present	1205005.68	8.87
	bed table bath	1036988.68	7.63
	sport leisure	988048.97	7.27
	computer accessories	911954.32	6.71
	Furniture Decoration	729762.49	5.37
	Cool Stuff	635290.85	4.67
	housewares	632248.66	4.65
	automotive	592720.11	4.36

```
# Merge order_items with products on product_id
order_items_products = df_order_items.merge(df_products[['product_id', 'product category']], on='product_id', how='inner')

# Clean product category
order_items_products['product category'] = [order_items_products['product category'].astype(str).strip()]
order_items_products = order_items_products[order_items_products['product category'] != '']


# Group by product category and sum the price
revenue_by_category = order_items_products.groupby('product category')['price'].sum()

# Calculate total revenue from joined data (same as SQL logic)
total_revenue = order_items_products['price'].sum()

# Calculate percentage
percentage_revenue_by_category = ((revenue_by_category / total_revenue) * 100).round(2)
percentage_revenue_by_category = percentage_revenue_by_category.sort_values(ascending=False)

# Show the top 10
print("Category Percentages (to match SQL):")
print(percentages_revenue_by_category.head(10))

# Plotting
plt.figure(figsize=(14, 7))
percentage_revenue_by_category.head(20).plot(kind='bar')
plt.title('Percentage of Total Revenue by Product Category (Top 20)')
plt.xlabel('Product Category')
plt.ylabel('Percentage of Total Revenue (%)')
plt.xticks(rotation=90)
plt.tight_layout()
plt.show()
```





- Identify the correlation between product price and the number of times a product has been purchased


```
WITH product_sales AS (
    SELECT
        oi.product_id,
        oi.price,
        COUNT(oi.order_id) AS purchase_count
    FROM
        order_items oi
    GROUP BY
        oi.product_id, oi.price
)
SELECT
    (
        COUNT(*) * SUM(price * purchase_count) - SUM(price) * SUM(purchase_count)
    ) /
    (
        SQRT(COUNT(*) * SUM(POWER(price, 2)) - POWER(SUM(price), 2)) *
        SQRT(COUNT(*) * SUM(POWER(purchase_count, 2)) - POWER(SUM(purchase_count), 2))
    ) AS correlation
FROM
    product_sales;
```

Result Grid	Filter Rows:
	correlation
▶	-0.0442186927706809

```
# Group by both product_id and price, count purchases
product_sales = (
    df_order_items
    .groupby(['product_id', 'price'])
    .agg(purchase_count=('order_id', 'count'))
    .reset_index())

# Calculate Pearson correlation (same formula as SQL)
correlation = product_sales['price'].corr(product_sales['purchase_count'])
print(f"Pearson correlation between product price and purchase count: {correlation}")

# Plot
plt.figure(figsize=(10, 6))
sns.scatterplot(x='price', y='purchase_count', data=product_sales, alpha=0.5)
plt.xlabel('Product Price')
plt.ylabel('Purchase Count')
plt.title('Product Price vs. Purchase Count')
plt.grid(True)
plt.show()
```





- Calculate the total revenue generated by each seller, and rank them by revenue

```
SELECT
    seller_id,
    total_revenue,
    RANK() OVER (ORDER BY total_revenue DESC) AS revenue_rank
from(SELECT
    order_items.seller_id,
    SUM(order_items.price) AS total_revenue
FROM
    order_items
group by order_items.seller_id)
as seller_revenue
order by revenue_rank;
```

Result Grid				Filter Rows:	Export:	Wrap Cell
	seller_id	total_revenue	revenue_rank			
▶	4869f7a5dfa277a7dca6462dcf3b52b2	229472.63	1			
	53243585a1d6dc2643021fd1853d8905	222776.05	2			
	4a3ca9315b744ce9f8e9374361493884	200472.92	3			
	fa1c13f2614d7b5c4749cbc52fecda94	194042.03	4			
	7c67e1448b00f6e969d365cea6b010ab	187923.89	5			

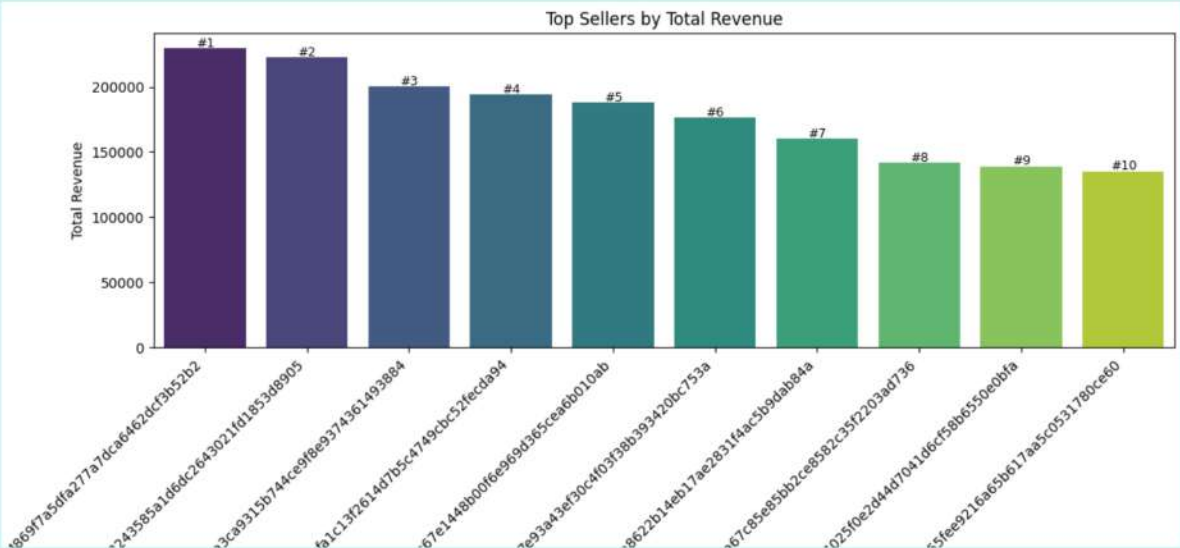
```
# Calculate total revenue per seller
seller_revenue = (df_order_items
    .groupby('seller_id', as_index=False)['price']
    .sum()
    .rename(columns={'price': 'total_revenue'}))

# Apply RANK
seller_revenue['revenue_rank'] = seller_revenue['total_revenue'].rank(
    method='dense', ascending=False).astype(int)

# Order by revenue_rank
seller_revenue = seller_revenue.sort_values('revenue_rank')
print(seller_revenue)

# Limit to top 10 or 20 for visualization (otherwise chart is too big)
top_sellers = seller_revenue.head(10)
plt.figure(figsize=(12, 6))
sns.barplot(x='seller_id', y='total_revenue', data=top_sellers, palette='viridis')
plt.title('Top Sellers by Total Revenue')
plt.xlabel('Seller ID')
plt.ylabel('Total Revenue')
plt.xticks(rotation=45, ha='right')
for i, row in top_sellers.reset_index(drop=True).iterrows():
    plt.text(i, row.total_revenue + 1000, f"#{row.revenue_rank}", ha='center', fontsize=9)
plt.tight_layout()
plt.show()
```

	seller_id	total_revenue	revenue_rank
857	4869f7a5dfa277a7dca6462dcf3b52b2	229472.63	1
1013	53243585a1d6dc2643021fd1853d8905	222776.05	2
881	4a3ca9315b744ce9f8e9374361493884	200472.92	3
3024	fa1c13f2614d7b5c4749cbc52fecda94	194042.03	4
1535	7c67e1448b00f6e969d365cea6b010ab	187923.89	5






# ADVANCED PROBLEMS


OBJECTIVE: GENERATE STRATEGIC AND CUSTOMER-CENTRIC INSIGHTS.

- Calculate the moving average of order values for each customer over their order history

```
SELECT
customer_id,
order_purchase_timestamp,
order_value,
AVG(order_value) OVER (
  PARTITION BY customer_id
  ORDER BY order_purchase_timestamp
  ROWS BETWEEN 2 PRECEDING AND CURRENT ROW
) AS moving_avg_3
FROM (
  SELECT
    o.customer_id,
    o.order_purchase_timestamp,
    SUM(oi.price) AS order_value
  FROM
    orders o
  JOIN
    order_items oi ON o.order_id = oi.order_id
  GROUP BY
    o.customer_id, o.order_purchase_timestamp, o.order_id
) AS order_history
ORDER BY
  customer_id, order_purchase_timestamp;
```


Result Grid






Filter Rows:

Export:



Wrap Cell Content:



	customer_id	order_purchase_timestamp	order_value	moving_avg_3
▶	00012a2ce6f8dcda20d059ce98491703	2017-11-14 16:08:00	89.80	89.800000
	000161a058600d5901f007fab4c27140	2017-07-16 09:40:00	54.90	54.900000
	0001fd6190edaaf884bcaf3d49edf079	2017-02-28 11:06:00	179.99	179.990000
	0002414f95344307404f0ace7a26f1d5	2017-08-16 13:09:00	149.90	149.900000
	000379cdec625522490c315e70c7a9fb	2018-04-02 13:42:00	93.00	93.000000
	0004164d20a9e969af783496f3408652	2017-04-12 08:35:00	59.99	59.990000
	000419c5494106c306a97b5635748086	2018-03-02 17:47:00	34.30	34.300000
	00046a560d407e99b969756e0b10f282	2017-12-18 11:08:00	120.90	120.900000

```
#Calculate order_value for each order
order_value_per_order = df_order_items.groupby('order_id')['price'].sum().reset_index(name='order_value')
order_history = pd.merge(
    df_orders[['order_id', 'customer_id', 'order_purchase_timestamp']],
    order_value_per_order, on='order_id', how='inner')
order_history['order_purchase_timestamp'] = pd.to_datetime(order_history['order_purchase_timestamp'])
order_history = order_history.sort_values(['customer_id', 'order_purchase_timestamp'])

# Calculate the moving average
order_history['moving_avg_3'] = (order_history.groupby('customer_id')['order_value'].rolling(window=3, min_periods=1).mean().reset_index(level=0, drop=True))

# Display the first 10 rows of the table
top_10_orders = order_history[['customer_id', 'order_purchase_timestamp', 'order_value', 'moving_avg_3']].head(10)
print(top_10_orders.to_string(index=False))

#Plot
plt.figure(figsize=(12, 6))
sns.lineplot(data=order_history, x='order_purchase_timestamp', y='moving_avg_3', color='red', linestyle='--', label="Moving Average")
plt.xlabel('Order Purchase Timestamp')
plt.ylabel('Order Value')
plt.title('Order Values and Moving Average by Customer')
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
plt.show()
```

	customer_id	order_purchase_timestamp	order_value	moving_avg_3
68055	00012a2ce6f8dcda20d059ce98491703	2017-11-14 16:08:00	89.80	89.80
9936	000161a058600d5901f007fab4c27140	2017-07-16 09:40:00	54.90	54.90
65380	0001fd6190edaaf884bcaf3d49edf079	2017-02-28 11:06:00	179.99	179.99
42834	0002414f95344307404f0ace7a26f1d5	2017-08-16 13:09:00	149.90	149.90
5843	000379cdec625522490c315e70c7a9fb	2018-04-02 13:42:00	93.00	93.00
73081	0004164d20a9e969af783496f3408652	2017-04-12 08:35:00	59.99	59.99
45797	000419c5494106c306a97b5635748086	2018-03-02 17:47:00	34.30	34.30
59517	00046a560d407e99b969756e0b10f282	2017-12-18 11:08:00	120.90	120.90



- CALCULATE THE CUMULATIVE SALES PER MONTH FOR EACH YEAR

```
select
year(order_purchase_timestamp),
month(order_purchase_timestamp),
sum(order_value) as monthly_sales,
sum(sum(order_value))over(order by year(order_purchase_timestamp), month(order_purchase_timestamp)) as cummalat
from (select
o.order_id,
o.order_purchase_timestamp,
sum(oi.price) as order_value
from orders o
join order_items oi on o.order_id = oi.order_id
group by o.order_id,o.order_purchase_timestamp) as order_values
group by
year(order_purchase_timestamp),
month(order_purchase_timestamp)
ORDER BY
YEAR(order_purchase_timestamp), MONTH(order_purchase_timestamp);
```

Result Grid				
Filter Rows:				
Export:				
Wrap Cell Content:				
	year(order_purchase_timestamp)	month(order_purchase_timestamp)	monthly_sales	cummulative_sales
►	2016	9	267.36	267.36
	2016	10	49507.66	49775.02
	2016	12	10.90	49785.92
	2017	1	120312.87	170098.79
	2017	2	247303.02	417401.81
	2017	3	374344.30	791746.11
	2017	4	359927.23	1151673.34
	2017	5	506071.14	1657744.48
	2017	6	433038.60	2090783.08

```
#Merge order_items with orders to get order_purchase_timestamp with every item
merged = pd.merge(df_order_items, df_orders[['order_id', 'order_purchase_timestamp']], on='order_id', how='inner')
merged['order_purchase_timestamp'] = pd.to_datetime(merged['order_purchase_timestamp'])

# Calculate order_value for each order (SUM(oi.price) per order)
order_value_df = merged.groupby(['order_id', 'order_purchase_timestamp'])['price'].sum().reset_index(name='order_value')
order_value_df['year'] = order_value_df['order_purchase_timestamp'].dt.year
order_value_df['month'] = order_value_df['order_purchase_timestamp'].dt.month

# Group by year and month, and sum order_value to get monthly_sales
monthly_sales = (
    order_value_df
    .groupby(['year', 'month'])['order_value']
    .sum()
    .reset_index(name='monthly_sales')
    .sort_values(['year', 'month']))

#Calculate cumulative_sales like the SQL window function
monthly_sales['cumulative_sales'] = monthly_sales['monthly_sales'].cumsum()
print(monthly_sales)

#Plot
plt.figure(figsize=(14, 7))
sns.lineplot(data=monthly_sales, x='month', y='cumulative_sales', hue='year', palette='viridis')
plt.xlabel('Month')
plt.ylabel('Cumulative Sales')
plt.title('Cumulative Sales Over Time (by Month and Year)')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

	year	month	monthly_sales	cumulative_sales
0	2016	9	267.36	267.36
1	2016	10	49507.66	49775.02
2	2016	12	10.90	49785.92
3	2017	1	120312.87	170098.79
4	2017	2	247303.02	417401.81
5	2017	3	374344.30	791746.11
6	2017	4	359927.23	1151673.34
7	2017	5	506071.14	1657744.48
8	2017	6	433038.60	2090783.08
9	2017	7	498031.48	2588814.56



• CALCULATE THE YEAR-OVER-YEAR GROWTH RATE OF TOTAL SALES

```
CREATE TEMPORARY TABLE order_sales AS
SELECT
  o.order_id,
  SUM(oi.price) AS order_total,
  o.order_purchase_timestamp
FROM
  orders o
  INNER JOIN order_items oi ON o.order_id = oi.order_id
WHERE
  o.order_purchase_timestamp IS NOT NULL
GROUP BY
  o.order_id, o.order_purchase_timestamp;

SELECT
  year,
  total_sales,
  LAG(total_sales) OVER (ORDER BY year) AS previous_year_sales,
  ROUND(
    CASE
      WHEN LAG(total_sales) OVER (ORDER BY year) IS NULL THEN 0
      ELSE ((total_sales - LAG(total_sales) OVER (ORDER BY year)) / LAG(total_sales) OVER (ORDER BY year))
    END
    ,2) AS YoY_Growth_Rate_Percent
FROM (
  SELECT
    YEAR(order_purchase_timestamp) AS year,
    SUM(order_total) AS total_sales
  FROM
    order_sales
  GROUP BY
    YEAR(order_purchase_timestamp)
) yearly
ORDER BY year;
```

	year	total_sales	previous_year_sales	YoY_Growth_Rate_Percent
▶	2016	49785.92	NULL	0.00
	2017	6155806.98	49785.92	12264.55
	2018	7386050.80	6155806.98	19.99

```
df_orders['order_purchase_timestamp'] = pd.to_datetime(df_orders['order_purchase_timestamp'], errors='coerce')
df_orders = df_orders.dropna(subset=['order_purchase_timestamp'])

# Merge orders with order_items to get price information per order
orders_with_prices = pd.merge(df_orders, df_order_items[['order_id', 'price']], on='order_id', how='inner')


# Calculate total sales per order
order_sales = orders_with_prices.groupby('order_id')['price'].sum().reset_index(name='order_total')
order_sales = pd.merge(order_sales, df_orders[['order_id', 'order_purchase_timestamp']], on='order_id', how='left')
order_sales['year'] = order_sales['order_purchase_timestamp'].dt.year

# Calculate total sales per year
yearly_sales = order_sales.groupby('year')['order_total'].sum().reset_index(name='total_sales')
yearly_sales = yearly_sales.sort_values('year')
yearly_sales['previous_year_sales'] = yearly_sales['total_sales'].shift(1)

# Calculate growth rate: (Current Year Sales - Previous Year Sales) / Previous Year Sales * 100
yearly_sales['YoY_Growth_Rate (%)'] = ((yearly_sales['total_sales'] - yearly_sales['previous_year_sales']) / yearly_sales['previous_year_sales']) * 100
yearly_sales['YoY_Growth_Rate (%)'] = yearly_sales['YoY_Growth_Rate (%)'].fillna(0)

print("Year-over-Year Growth Rate of Total Sales:")
print(yearly_sales)

# Plotting the YoY Growth Rate
plt.figure(figsize=(10, 6))
sns.lineplot(x='year', y='YoY_Growth_Rate (%)', data=yearly_sales, marker='o')
plt.xlabel('Year')
plt.ylabel('YoY Growth Rate (%)')
plt.title('Year-over-Year Growth Rate of Total Sales')
plt.grid(True)
plt.show()
```





- **Calculate the retention rate of customers, defined as the percentage of customers who make another purchase within 6 months of their first purchase**

```
-- Calculate the retention rate of customers, defined as the percentage of customers who make another
SELECT order_purchase_timestamp FROM orders LIMIT 5;
WITH first_purchase AS (
    SELECT
        c.customer_unique_id,
        MIN(o.order_purchase_timestamp) AS first_order_date
    FROM
        orders o
    JOIN
        customers c ON o.customer_id = c.customer_id
    GROUP BY
        c.customer_unique_id
),
retained_customers AS (
    SELECT DISTINCT c.customer_unique_id
    FROM orders o
    JOIN customers c ON o.customer_id = c.customer_id
    JOIN first_purchase fp ON c.customer_unique_id = fp.customer_unique_id
    WHERE
        o.order_purchase_timestamp > fp.first_order_date
        AND o.order_purchase_timestamp <= DATE_ADD(fp.first_order_date, INTERVAL 6 MONTH)
)
SELECT
    (SELECT COUNT(*) FROM retained_customers) AS retained_customers,
    (SELECT COUNT(*) FROM first_purchase) AS total_customers,
    ROUND(100.0 * (SELECT COUNT(*) FROM retained_customers) / (SELECT COUNT(*) FROM first_purchase), 2)
```

Result Grid	Filter Rows:	Export:	Wrap Cell Content:
	retained_customers	total_customers	retention_rate_percent
▶	1804	96096	1.88

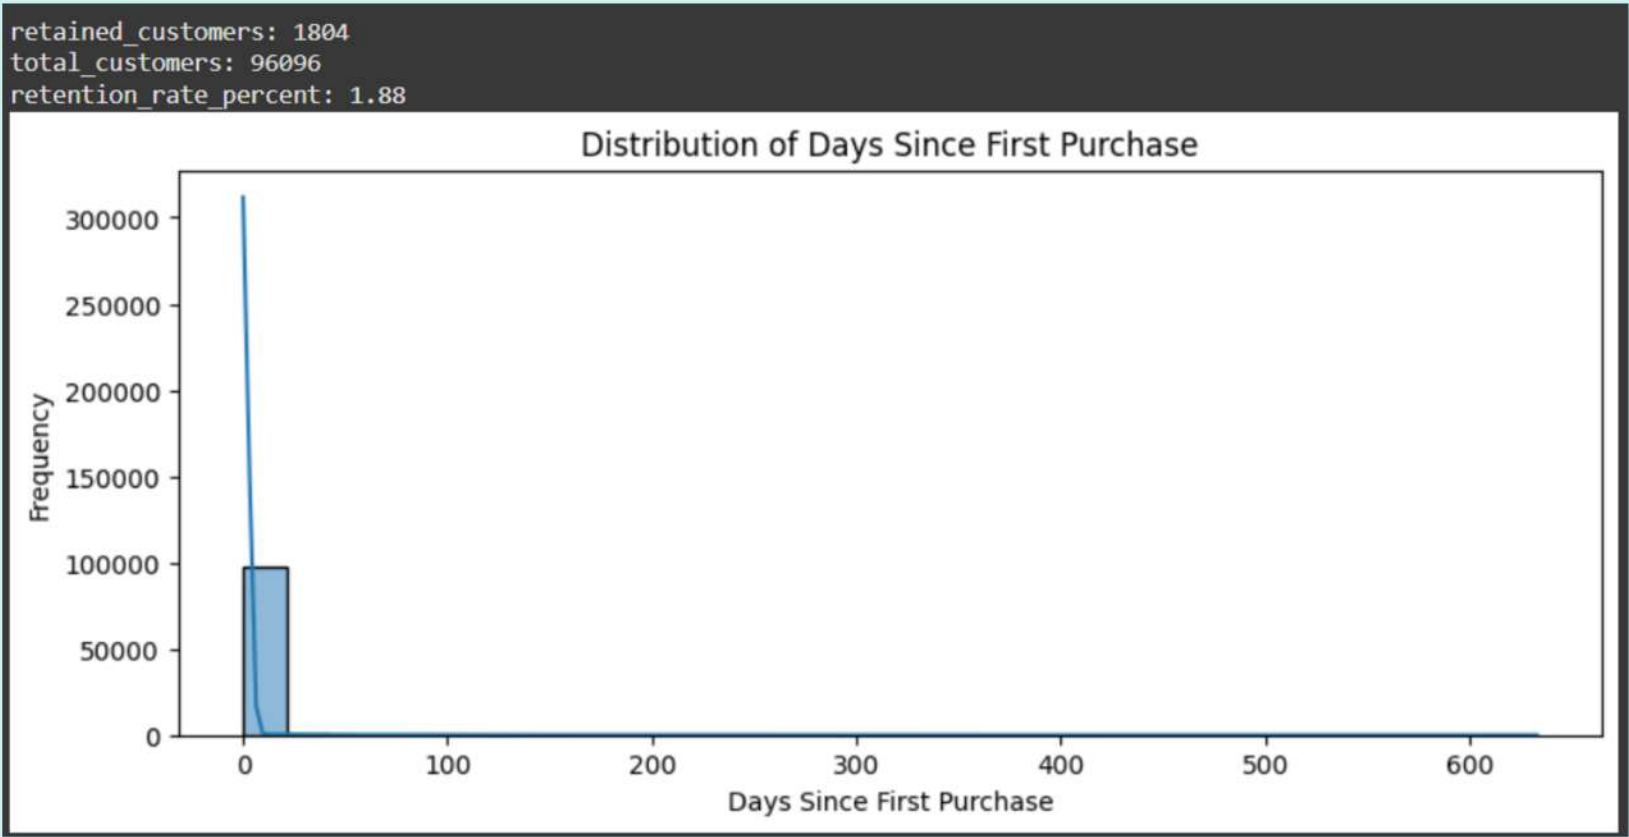
```
orders = pd.merge(df_orders, df_customers[['customer_id', 'customer_unique_id']], on='customer_id')
orders['order_purchase_timestamp'] = pd.to_datetime(orders['order_purchase_timestamp'], format='%d-%m-%Y %H:%M', errors='coerce')
orders.dropna(subset=['order_purchase_timestamp'], inplace=True)

# Find first purchase
first_order_timestamp = (orders.groupby('customer_unique_id')['order_purchase_timestamp'].min().reset_index()
    .rename(columns={'order_purchase_timestamp': 'first_order_timestamp'}))
orders = pd.merge(orders, first_order_timestamp, on='customer_unique_id', how='left')

# Calculate the date exactly 6 months after the first purchase timestamp
orders['six_months_later'] = orders['first_order_timestamp'] + pd.DateOffset(months=6)
retained_orders_within_window = orders[(orders['order_purchase_timestamp'] > orders['first_order_timestamp']) &
    (orders['order_purchase_timestamp'] <= orders['six_months_later'])]

# Count the number of unique customers
retained_customers = retained_orders_within_window['customer_unique_id'].nunique()
total_customers = orders['customer_unique_id'].nunique()
retention_rate = round(100 * retained_customers / total_customers, 2)
print('retained_customers:', retained_customers)
print('total_customers:', total_customers)
print('retention_rate_percent:', retention_rate)
orders['days_since_first'] = (orders['order_purchase_timestamp'] - orders['first_order_timestamp']).dt.days

# Plot
plt.figure(figsize=(10, 4))
sns.histplot(orders['days_since_first'], bins=30, kde=True)
plt.xlabel('Days Since First Purchase')
plt.ylabel('Frequency')
plt.title('Distribution of Days Since First Purchase')
plt.show()
```





- Identify the top 3 customers who spent the most money in each year

```
SELECT *
FROM (
    SELECT
        YEAR(o.order_purchase_timestamp) AS year,
        c.customer_unique_id,
        SUM(p.payment_value) AS total_spent,
        ROW_NUMBER() OVER (PARTITION BY YEAR(o.order_purchase_timestamp) ORDER BY SUM(p.payment_value) DESC)
    FROM orders o
    JOIN customers c ON o.customer_id = c.customer_id
    JOIN payments p ON o.order_id = p.order_id
    GROUP BY year, c.customer_unique_id
) ranked
WHERE rn <= 3
ORDER BY year, total_spent DESC;
```

Result Grid					Filter Rows:	Export:	Wrap Cell
	year	customer_unique_id	total_spent	rn			
▶	2016	fdaa290acb9eeacb66fa7f979baa6803	1423.55	1			
	2016	753bc5d6efa9e49a03e34cf521a9e124	1400.74	2			
	2016	b92a2e5e8a6eabcc80882c7d68b2c70b	1227.78	3			
	2017	0a0a92112bd4c708ca5fde585afaa872	13664.08	1			
	2017	da122df9eeddfedc1dc1f5349a1a690c	7571.63	2			
	2017	dc4802a71eae9be1dd28f5d788ceb526	6929.31	3			
	2018	46450c74a0d8c5ca9395da1daac6c120	9553.02	1			
	2018	763c8b1c9c68a0229c42c9fc6f662b93	7274.88	2			
	2018	459bef486812aa25204be022145caa62	6922.21	3			

```
merged = pd.merge(df_orders, df_customers[['customer_id', 'customer_unique_id']], on='customer_id', how='inner')
merged = pd.merge(merged, df_payments[['order_id', 'payment_value']], on='order_id', how='inner')

# Convert order_purchase_timestamp to datetime with the correct format
merged['order_purchase_timestamp'] = pd.to_datetime(merged['order_purchase_timestamp'],
                                                    format='%d-%m-%Y %H:%M', errors='coerce')

merged.dropna(subset=['order_purchase_timestamp'], inplace=True)
merged['year'] = merged['order_purchase_timestamp'].dt.year

# Group by year and customer, sum payment_value
customer_yearly_spend = (
    merged.groupby(['year', 'customer_unique_id'])['payment_value']
    .sum()
    .reset_index(name='total_spent')
)
customer_yearly_spend = customer_yearly_spend.sort_values(['year', 'total_spent'], ascending=[True, False])
top3_each_year = customer_yearly_spend.groupby('year').head(3).reset_index(drop=True)
print(top3_each_year)
```

	year	customer_unique_id	total_spent
0	2016	fdaa290acb9eeacb66fa7f979baa6803	1423.55
1	2016	753bc5d6efa9e49a03e34cf521a9e124	1400.74
2	2016	b92a2e5e8a6eabcc80882c7d68b2c70b	1227.78
3	2017	0a0a92112bd4c708ca5fde585afaa872	13664.08
4	2017	da122df9eeddfedc1dc1f5349a1a690c	7571.63
5	2017	dc4802a71eae9be1dd28f5d788ceb526	6929.31
6	2018	46450c74a0d8c5ca9395da1daac6c120	9553.02
7	2018	763c8b1c9c68a0229c42c9fc6f662b93	7274.88
8	2018	459bef486812aa25204be022145caa62	6922.21

# Key Insights

## **CUSTOMER RETENTION IS LOW:**

Only about 1.8% of customers made a repeat purchase within 6 months of their first order, indicating a potential area for growth in customer engagement and loyalty programs.

## **SALES SHOW YEAR-OVER-YEAR GROWTH:**

The dataset reveals a consistent increase in total sales year over year, demonstrating healthy business expansion and a growing customer base.

## **TOP CUSTOMERS DRIVE REVENUE:**

The top 3 customers in each year accounted for a disproportionately large share of annual revenue, emphasizing the value of nurturing high-spending customers.

## **MONTHLY SALES FLUCTUATIONS:**

Certain months saw spikes in sales, possibly due to seasonality or marketing campaigns. Identifying these peaks can help with inventory and marketing planning.

## **ORDER VALUE TRENDS:**

The moving average analysis showed that while most customers place orders of moderate value, a few high-value purchases significantly boost overall revenue.



# Challenges

## **Data Consistency:**

Merging multiple CSV files required careful handling of inconsistent columns, missing values, and duplicate entries, especially when joining orders, payments, and customer data.

## **SQL vs Python Logic Differences:**

Ensuring that calculations (such as moving averages and window functions) produced identical results in both SQL and Python was sometimes tricky due to subtle differences in how each tool handles dates and grouping.

## **Handling Time Data:**

Some discrepancies in retention rates and time-based calculations arose due to variations in how time intervals (like “6 months”) are treated in SQL versus Python.

## **Performance:**

Processing large datasets with groupby and merge operations requires optimization, especially in Python, to avoid long execution times.

# Conclusions

## **DATA-DRIVEN STRATEGY:**

The analysis provided actionable insights into sales patterns, customer retention, and the impact of top customers on business growth.

## **GROWTH OPPORTUNITIES:**

Improving customer retention even by a few percentage points could result in substantial increases in revenue.

## **BUSINESS FOCUS:**

Efforts should focus on engaging high-value customers and converting first-time buyers into repeat purchasers through targeted campaigns.

## **SQL + PYTHON SYNERGY:**

Combining SQL for data extraction with Python for visualization and deeper analysis enabled a comprehensive exploration of the dataset.



# Thank you very much!

For More Info-

Email – [mohi.workspace@gmail.com](mailto:mohi.workspace@gmail.com)

LinkedIn Id- [www.linkedin.com/in/mohi-gupta-17b8b51b6](https://www.linkedin.com/in/mohi-gupta-17b8b51b6)

Github-<https://github.com/MohiGupta26>